




Fast and Memory-Efficient Tucker Decomposition for Answering Diverse Time Range Queries

Jun-Gi Jang and U Kang
Data Mining Lab
Dept. of CSE
Seoul National University

Overview

- Q. Given a temporal dense tensor and a time range (e.g., January -March 2019), how can we efficiently analyze the tensor in the given time range?
- A. **Zoom-Tucker** enables us to analyze the tensor in the given range, quickly and memory-efficiently




Outline

→ ■ Introduction

- Proposed Method
- Experiments
- Conclusion

Temporal Dense Tensors

- Several real-world data are represented as **temporal dense tensors**
 - One dimension corresponds to time
 - Most entries of a tensor are measured




Stocks

3-way tensor

Index: (stock, feature, time)

Value: measurement




Video

3-way tensor

Index: (width, height, time)

Value: measurement



Traffic Volume

3-way tensor

Index: (sensor, frequency, time)



Value: measurement

Tucker Decomposition

- Given an N -order tensor $\mathcal{X} \in \mathbb{R}^{I_1 \times \dots \times I_N}$, rank J_1, \dots, J_N
- Obtain factor matrices $\mathbf{A}^{(n)} \in \mathbb{R}^{I_n \times J}$ for $n = 1, \dots, N$ and core tensor $\mathcal{G} \in \mathbb{R}^{J_1 \times \dots \times J_N}$
- Objective function

$$\min_{\mathcal{G}, \{\mathbf{A}^{(k)}\}_{k=1}^N} \|\mathbf{X}_{(n)} - \mathbf{A}^{(n)} \mathbf{G}_{(n)} (\otimes_{k \neq n}^N \mathbf{A}^{(k)T})\|_F^2$$

3-order tensor



\otimes : Kronecker product
 $(\otimes_{k \neq n}^N \mathbf{A}^{(k)T})$: the entire Kronecker product of $\mathbf{A}^{(k)T}$ in descending order for $k = N, \dots, n+1, n-1, \dots, 1$
 $\mathbf{X}_{(n)}$: mode- n matricization of \mathcal{X}
 $\mathbf{G}_{(n)}$: mode- n matricization of \mathcal{G}



- ALS (Alternating Least Square) approach

- Iteratively updates a factor matrix of a mode while fixing all factor matrices of other modes

Applications

■ Several applications for Tucker decomposition

- Dimensionality reduction, concept discovery, trend analysis, anomaly detection, and clustering



Dimensionality Reduction

Lossy compression for a temporal dense tensor

Sensor type 1,3,4,7



Sensor type 2,5,6

Trend Analysis


Find main trends using latent factors of the time dimension

Clustering

Find similar objects using latent factors

Time Range Query

- Several users are interested in investigating patterns of **diverse** time ranges using Tucker decomposition



Problem Definition

Problem: time range query problem on temporal dense tensor

■ Given

- A temporal dense tensor $\mathcal{X} \in \mathbb{R}^{I_1 \times \dots \times I_{N-1} \times I_N}$
 - Assume the last dimension is the time dimension
- A time range query $[t_s, t_e]$

■ Find

- The Tucker results of the sub-tensor of \mathcal{X} included in the given range $[t_s, t_e]$
 - The Tucker results include factor matrices $\tilde{\mathbf{A}}^{(1)}, \dots, \tilde{\mathbf{A}}^{(N)}$, and core tensor $\tilde{\mathbf{G}}$

Challenges

- **(Limitation)** Previous works are tailored for performing Tucker decomposition of only the whole tensor once
 - For each time range, performing Tucker decomposition from scratch requires **high time and space costs**
 - A few methods with a preprocessing phase are still **unsatisfactory in terms of time, space, and accuracy** on the problem
 - Before time range is given, they preprocess a given tensor, and perform Tucker decomposition with the preprocessed tensor for each time range query
- Need to address the following challenges
 1. Deal with diverse time ranges
 2. Minimize computational costs
 3. Avoid huge intermediate data

Outline


- Introduction
- ➡ ■ Proposed Method
- Experiments
- Conclusion

Proposed Method - Overview

- We propose Zoom-Tucker (Zoomable Tucker decomposition)
 - A **fast** and **memory-efficient** Tucker decomposition method for diverse time range queries
- Zoom-Tucker consists of the two phases
 - **Preprocessing phase**
 - Compresses a given temporal tensor block by block before time range queries are given
 - **Query phase**
 - Answers a given time range query by exploiting compression results obtained in the preprocessing phase

Preprocessing Phase

- Before time range queries are given, Zoom-Tucker compresses a given temporal tensor block by block along the time dimension
 - Split a given temporal tensor into sub-tensors along the time dimension
 - For each block tensor, we perform Tucker decomposition




- The advantages of the preprocessing phase
 1. Generate **small results** compared to an input tensor (support high efficiency of the query phase)
 2. Capture **local temporal information** (reduce error increase)

Query Phase

■ Goal

- Given a time range query, efficiently perform Tucker decomposition of the sub-tensor included in the time range



- A naïve approach would reconstruct the preprocessed results in a given range, and then compute Tucker decomposition
 - Unsatisfactory in terms of time and space costs

$\tilde{\mathbf{X}}_{(n)}$: the mode- n matricized version of a sub-tensor for a time range

$\tilde{\mathbf{A}}^{(n)}$: the n -th factor matrix

$\tilde{\mathbf{G}}_{(n)}$: the core tensor

■ Optimization problem

$$\min_{\tilde{\mathbf{G}}, \{\tilde{\mathbf{A}}^{(k)}\}_{k=1}^N} \left\| \tilde{\mathbf{X}}_{(n)} - \tilde{\mathbf{A}}^{(n)} \tilde{\mathbf{G}}_{(n)} \left(\bigotimes_{k \neq n}^N \tilde{\mathbf{A}}^{(k)T} \right) \right\|_F^2$$


■ Procedure

1. Given a time range $[t_s, t_e]$, load Tucker results included in the time range
2. Adjust the first and the last blocks to fit the range
3. **Alternatively update factor matrices and core tensor using the Tucker results**
4. Repeatedly performs step 3 until convergence

■ Update rule


$$\begin{aligned} \tilde{\mathbf{A}}^{(n)} &\leftarrow \tilde{\mathbf{X}}_{(n)} \left(\bigotimes_{k \neq n}^N \tilde{\mathbf{A}}^{(k)} \right) \tilde{\mathbf{G}}_{(n)}^T \left(\tilde{\mathbf{G}}_{(n)} \left(\bigotimes_{k \neq n}^N \tilde{\mathbf{A}}^{(k)T} \tilde{\mathbf{A}}^{(k)} \right) \tilde{\mathbf{G}}_{(n)}^T \right)^{-1} \\ &= \boxed{\tilde{\mathbf{X}}_{(n)} \left(\bigotimes_{k \neq n}^N \tilde{\mathbf{A}}^{(k)} \right) \tilde{\mathbf{G}}_{(n)}^T (\mathbf{C}^{(n)})^{-1}} \end{aligned}$$

Efficiently compute this term



Main Ideas


- Main ideas for the efficient query phase
 1. In our update rule, replace a sub-tensor \tilde{X} of a time range query with **preprocessed block Tucker results**
 2. **Decouple** block Tucker results in a time range query, obtained in the preprocessing phase
 3. Carefully determining the **order** of computations



- With the above two techniques, we **significantly reduce** the intermediate data and computational costs for a given query

Non-temporal Mode

- By exploiting the two techniques, we decouple block results along the time dimension
- Compute each block and then sum up the results



- In $(\mathbf{B}^{<i>})^{(n)} \in \mathbb{R}^{J \times J}$ and $\mathbf{C}^{(n)} \in \mathbb{R}^{J \times J}$, there are the following computations (**low time and space costs**)
 - $\mathbf{U}^T \mathbf{V} \in \mathbb{R}^{J \times J}$ ($\mathbf{U}, \mathbf{V} \in \mathbb{R}^{I \times J}$)
 - n -mode products between a core tensor and $\mathbf{U}^T \mathbf{V}$
 - Matrix multiplication between mode- n matricization of two core tensors

Temporal Mode & Core Tensor

- Decouple block results along the time dimension
- Compute each block

$$\tilde{\mathbf{A}}^{(N)} \leftarrow \begin{bmatrix} (\mathbf{A}^{<S>})^{(N)} (\mathbf{B}^{<S>})^{(N)} \\ \vdots \\ (\mathbf{A}^{<E>})^{(N)} (\mathbf{B}^{<E>})^{(N)} \end{bmatrix} (\mathbf{C}^{(N)})^{-1}$$

- $(\mathbf{B}^{<i>})^{(N)} \in \mathbb{R}^{J \times J}$ and $\mathbf{C}^{(N)} \in \mathbb{R}^{J \times J}$ are also computed with low computational costs
- At the end of each iteration, the core tensor is also updated by exploiting the main ideas

With the main ideas, the dominant cost to update a factor matrix is mainly proportional to **the size I** of dimension and **the number B** of blocks \Rightarrow **fast** and **memory-efficient**

Outline

- Introduction
- Proposed Method
- ➡ ■ **Experiments**
- Conclusion

Experimental Questions

- Answer the following questions:

- Q1. Performance trade-off.** Does Zoom-Tucker provide the best trade-off between query time and reconstruction error?
- Q2. Space cost.** What is the space cost of Zoom-Tucker and competitors for preprocessed results?
- Q3. Effects of the block size b .** How does a block size b affect query time and reconstruction error of Zoom-Tucker?
- Q4. Discovery.** What pattern does Zoom-Tucker discover in different time ranges?

Real-world Datasets

- Use 6 real-world datasets

Dataset	Dimensionality	Length $l_{[t_s, t_e]}$ of Time Range	Summary
Boats ¹ [37]	$320 \times 240 \times 7000$	(128, 2048)	Video
Walking Video [25]	$1080 \times 1980 \times 2400$	(128, 2048)	Video
Stock ³	$3028 \times 54 \times 3050$	(128, 2048)	Time series
Traffic ⁴ [30]	$1084 \times 96 \times 2000$	(64, 1024)	Traffic volume
FMA ⁵ [8]	$7994 \times 1025 \times 700$	(32, 512)	Music
Absorb ⁶	$192 \times 288 \times 30 \times 1200$	(64, 1024)	Climate

- The last dimension is the time dimension
- Length of time range
 - For each dataset, we use two kinds of time range queries: narrow and wide time ranges

Experimental Setting

- Machine
 - A workstation with a single CPU (Intel Xeon E5-2630 v4 @ 2.2GHz), and 512GB memory
- Target Rank
 - 10
- Block size b
 - 50
- Reconstruction error

$$\frac{\|\mathcal{X} - \widehat{\mathcal{X}}\|_F^2}{\|\mathcal{X}\|_F^2}$$


- \mathcal{X} is an input tensor
- $\widehat{\mathcal{X}}$ is the tensor reconstructed from factor matrices and core tensor

Competitor

- Compare Zoom-Tucker with the following Tucker decomposition approaches
 - **D-Tucker**
 - A SVD based Tucker decomposition method
 - **Tucker-ts and Tucker-ttmts**
 - A sketching based Tucker decomposition method
 - **MACH**
 - A sampling based Tucker decomposition method
 - **Tucker-ALS**
 - An implementation of ALS algorithm in Tensor Toolbox
 - **RTD**
 - A Tucker decomposition method with a randomized algorithm

Q1. Performance Trade-off


- The lower-left region indicates better performance



Zoom-Tucker **outperforms** the competitors based on Tucker-ALS for both narrow and wide time ranges while having comparable errors

Q2. Space Cost

- Compare space cost for storing preprocessed results



- *Input Tensor* corresponds to the space cost of Tucker-ALS, Tuckerts, Tucker-ttmts, and RTD
- Zoom-Tucker requires up to **230× less space** than competitors

Q3. Effect of block size


- Measure the running time and error with respect to block size b

- Narrow time range queries**

- Trade-off relationships
 - $b \uparrow$
 - Running time \downarrow
 - Error \uparrow

- Wide time range queries**

- The running time is **inversely proportional** to b
 - The reconstruction error is **not sensitive** to b




Q4. Discovery

- Given a Korean stock dataset in the form of (stock, features, time), analyze the trend change by comparing results of two time range queries
 - Analyze the change of yearly trend of *Samsung Electronics* in the years **2013 (Query 1)** and **2018 (Query 2)**
- 1. Perform Zoom-Tucker for the two time range queries and get the factor matrix $\tilde{\mathbf{A}}^{(1)}$ each of whose rows contain the latent features of a stock
- 2. Manually pick 33 smartphone-related stocks and 46 semiconductor-related stocks
- 3. For each query, compare the cosine distance between the latent feature vectors of each stock and *Samsung Electronics*

Q4. Discovery

- Analyze trend change by comparing results of the two time range queries



- A clear change of the distances between 2013 and 2018
 - Samsung Electronics is more close to smartphone-related stocks in 2013
 - but to semiconductor-related stocks in 2018
- This result exactly reflects the sales trend of Samsung Electronics
 - In 2013, the annual sales of its smartphone division ↑
 - In 2018, those of its semiconductor division ↑

Zoom-Tucker enables us to efficiently explore diverse time ranges

Outline

- Introduction
- Proposed Method
- Experiments
- ➡ ■ Conclusion

Conclusion

- Zoom-Tucker answers diverse time range queries on a dense temporal tensor **quickly** and **memory-efficiently**
 - Compress a given temporal tensor block by block along the time dimension
 - Perform Tucker decomposition by elaborately using compression results, every time a time range is given
- Zoom-Tucker **outperforms** the previous Tucker decomposition methods based on ALS
- Zoom-Tucker provides opportunities to extract **unknown and interesting patterns** in diverse time ranges



Thank you !

<https://datalab.snu.ac.kr/zoomtucker>